

# Supporting Information

## Topological Systems with CP2K

Thomas D. Kühne, Alexander Cerjan, Vladislav Efremkin,  
Hermann Schulz-Baldes, and Emil Prodan

Corresponding author: [tkuehne@cp2k.org](mailto:tkuehne@cp2k.org)

### S1 Scope and Reproduction Boundaries

This working document combines the retained band-topology calculations transferred from the k-point manuscript with separately validated transport, localizer, and quadratic-pseudospectrum developments. The original results are not recomputed or relabelled by moving the manuscript text. The development status is recorded as of 25 September 2026. A source freeze and complete public reproduction archive remain submission requirements.

#### Applications versus test calculations

TopoHSE-DB and the submitted generative-predictive discovery study are cited in the main text as materials-science applications. Their VASP-based calculations are not part of the CP2K validation records collected here. In particular, their database labels and candidate counts do not validate the new CP2K localizer or transport implementations.

The smaller calculations below are tests and demonstrations. The companion k-point paper additionally reports 18 converged Si/Al calculations in conventional cubic cells with eight Si atoms ( $a = 5.43$  Å) and four Al atoms ( $a = 4.05$  Å). These use PBE, DZVP-MOLOPT-SR-GTH bases, GTH-PBE pseudopotentials, 600/60 Ry grid cutoffs, and an SCF threshold of  $10^{-9}$ . Al uses Fermi-Dirac occupations at 1000 K, so its compared quantity is the electronic free energy. The same-mesh comparison uses  $4 \times 4 \times 4$  sampling: 64 full-grid points, 32 with time reversal, and four with either K290 or SPGLIB atomic symmetry. The maximum cell-energy/free-energy difference is  $1.4 \times 10^{-12}$  hartree; the Al atomic-symmetry stress difference is 0.098 bar. These tests demonstrate representation consistency, not independently converged material properties or general force accuracy. Their source and raw records remain those of the companion manuscript; no new Si/Al calculations were performed for the present discussion.

### S2 Native Wilson-Loop Test and Demonstration Calculations

The topology examples form a separate calculation set from the Si/Al mainline tests. The native extension supplies directed Gaussian-basis overlaps, second-variational spinors, and the Wilson-loop analysis described in the main text. The retained calculations were completed on 17 September 2026 with one MPI process and one OpenMP thread. Source provenance and artifact checksums are recorded in `validation/topology/README.md` and the figure manifest.

Both examples use PBE, the DZVP-MOLOPT-GGA-GTH basis from `BASIS_MOLOPT_UZH`, SOC pseudopotentials from `GTH_SOC_POTENTIALS`, a 200 Ry density cutoff and a 40 Ry relative cutoff. The potential labels are `GTH-PBE-q8` for neon and `GTH-PBE-q4` for tin. Complex orbitals and explicit full SCF meshes are used. The neon reference contains one atom at fractional coordinate  $(1/2, 1/2, 1/2)$  in a cubic cell of side 5 Å, with a  $2 \times 2 \times 2$  SCF mesh and convergence threshold  $10^{-9}$ . The stanene cell has lattice vectors  $(4.674, 0, 0)$ ,  $(2.337, 4.047802737288466, 0)$ , and  $(0, 0, 20)$  Å, and two tin atoms at fractional positions  $(1/3, 1/3, 0.5213)$  and  $(2/3, 2/3, 0.4787)$ . It therefore

has a buckling height of 0.852 Å. These fixed geometric parameters follow the pristine-stanene reference reported by Zhang et al.<sup>1</sup> An  $8 \times 8 \times 1$  SCF mesh and convergence threshold  $10^{-8}$  are used for stanene. The examples test the topology implementation at these fixed settings, rather than constitute a new structural or band-gap prediction.

Eight occupied spinors are retained in each case. The stanene second-variational calculation includes all 26 scalar basis states, giving 52 spinor states before occupied-band selection. For neon, five scalar states produce ten spinors. The selected half-plane is  $\mathbf{k}(s, t) = (t, s/2, 0)$  in fractional reciprocal coordinates, with  $0 \leq s, t \leq 1$ . The loop closes with reciprocal shift  $(1, 0, 0)$ . Starting meshes contain four points on each of three loops for neon and 12 points on each of 13 loops for stanene. Refinement doubles the points per loop and inserts a midpoint between neighboring transverse loops. The accepted surfaces have  $8 \times 5$  and  $192 \times 193$  points, respectively. The last center changes are  $1.28 \times 10^{-5}$  and  $2.73 \times 10^{-4}$  in reduced units. The smallest link singular values are 0.99826 and 0.95248.

The native invariants are zero for neon and one for stanene. Z2Pack 2.2.1 gives the same stanene parity using either the raw overlaps or their unitary polar factors. For the latter, the maximum Wilson-eigenvalue difference from the native calculation is  $7.95 \times 10^{-14}$ . This comparison tests the loop algebra with identical input matrices. Separate operator checks enforce zero-step normalization, reverse-link adjoint relations, loop reversal, and atomic-position phases. A bonded two-atom full-spinor check gives maximum adjoint and zero-step residuals of  $2.79 \times 10^{-14}$  and  $2.00 \times 10^{-13}$ , respectively.

An additional adaptive Z2Pack-driven stanene surface uses 42 CP2K line calculations and ends with 15 accepted transverse lines, all satisfying the line-position, movement, and gap checks, and again gives  $\nu = 1$ . This calculation chooses its sampling independently but uses the same CP2K Hamiltonian and overlap implementation. The main-text figure shows the native samples without interpolation or artificial separation of degenerate centers. It is regenerated by `python3 validation/topology/plot_wilson.py`, which requires NumPy and Matplotlib. Raw `.wilson` files, inputs, logs, comparison reports, generated comma-separated figure data, and checksums accompany the figure.

## S3 Native Chern and Cross-Geometry Extension

### S3.1 Source versions and independent analysis routes

The Chern and cross-geometry checks were completed on 18 September 2026 with a separate development executable. They are distinct from the retained 17 September neon/stanene calculations and do not reassign those earlier results to a newer implementation. Source patches, executable records, library versions, and independent comparison scripts accompany the validation archive. The numerical comparison tools require NumPy and SciPy; Z2Pack 2.2.1 is an optional independent analysis dependency.

Table S1: Topology analysis routes and runtime requirements. All native routes act on Gaussian-basis electronic states without constructing maximally localized Wannier orbitals.

| Route                         | Quantity and sampling                                                                           | Dependencies and scope                                                                            |
|-------------------------------|-------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------|
| Native Wilson/ $\mathbb{Z}_2$ | Wilson spectra, Berry phases, and time-reversal half-plane $\mathbb{Z}_2$ .                     | CP2K native kernels, including second-variational SOC. No Z2Pack or Wannier90 runtime dependency. |
| Native Wilson/ $C_1$          | Determinant-phase winding on a full closed reciprocal surface, selected by <code>CHERN</code> . | CP2K native kernels. Fixed-rank isolated subspace, closure and refinement checks.                 |

<sup>1</sup>R.-w. Zhang et al., Room temperature quantum spin Hall insulator in ethynyl-derivative functionalized stanene films, *Sci. Rep.* **6**, 18879 (2016), <https://doi.org/10.1038/srep18879>.

| Route                          | Quantity and sampling                                                                                                            | Dependencies and scope                                                                                                                                     |
|--------------------------------|----------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Native inversion analysis      | Even/odd irreps, Fu–Kane indices and inversion-subgroup EBR signatures at TRIM.                                                  | Separate native extension, Section S4. No general space-group catalogue.                                                                                   |
| Native parameter-mesh analysis | $C_1$ on two parameters and the second Chern character pairing $C_2$ on four parameters, using physical cross-geometry overlaps. | <code>topology_phasons</code> reads <code>STATE_EXPORT</code> files. One process with optional OpenMP. No Python, Z2Pack, or Wannier90 runtime dependency. |
| External Z2Pack route          | Independent loop/surface analysis of CP2K overlaps, with optional adaptive CP2K sampling.                                        | Python and Z2Pack. Optional analysis and cross-validation, not the implementation of the native invariants.                                                |

### S3.2 Parameter-mesh discretization and physical closure

Let  $Q_\mu(x)$  be the polar-unitary selected-state overlap directed from mesh vertex  $x$  to  $x + \hat{\mu}$ . All boundary links include the specified physical sewing. The oriented plaquette is

$$P_{\mu\nu}(x) = Q_\mu(x)Q_\nu(x + \hat{\mu})Q_\mu^\dagger(x + \hat{\nu})Q_\nu^\dagger(x). \quad (\text{S1})$$

With this link orientation, the two-parameter estimator is

$$C_1^{\text{mesh}} = -\frac{1}{2\pi} \sum_x \text{Arg det } P_{12}(x). \quad (\text{S2})$$

Its sign matches the determinant Wilson-phase winding convention of the main text. For four parameters, define the full principal matrix logarithm  $F_{\mu\nu}(x) = \log P_{\mu\nu}(x)$ . Each  $F_{\mu\nu}$  is anti-Hermitian and already includes the plaquette area. The estimator is

$$C_2^{\text{mesh}} = -\frac{1}{4\pi^2} \sum_x \text{Re Tr } [F_{12}F_{34} - F_{13}F_{24} + F_{14}F_{23}](x). \quad (\text{S3})$$

The full matrices are needed for a non-Abelian occupied subspace. Equation (S3) approximates the second Chern character pairing and is not generally the second Chern class. Its finite-mesh value is reported without integer rounding.

Physical links use  $C_a^\dagger O_{ab} C_b$ , with both spinor components summed when present. The cross-geometry AO matrix  $O_{ab}$  retains the moving contracted Gaussian centers, periodic images, and reciprocal boundary phases. The cell, atom count and selected rank remain fixed throughout the family. Endpoint checks use explicit atom permutations, lattice translations, basis matching, and selected-subspace sewing. The code reports AO-metric errors, sampled selected/excluded-band gaps, minimum link singular values, and principal plaquette phases. A refined mesh must preserve the same physical family. Sampled isolation and invariant convergence do not by themselves establish a bulk spectral gap between samples. An indirect gap is interpreted only for a lowest-band prefix with a justified common energy reference across geometries.

### S3.3 Retained validation results

The extension was rebuilt with GNU Fortran 16 and OpenBLAS. Its curvature and snapshot unit tests and the existing Wilson/ $C_1/\mathbb{Z}_2$  unit test passed with two OpenMP threads. All four inputs in `QS/regtest-topology` passed with two MPI ranks and two OpenMP threads per rank. The auxiliary parameter-mesh executable itself was tested as one process, not as an MPI-distributed postprocessor.

The native tests cover singular or underresolved links, metric and seam failures, determinant phases, non-Abelian gauge covariance, and model refinement. For the occupied doublet of the

four-dimensional Dirac model at mass  $-3$ , the  $8^4$  estimate is  $-0.7424128463$  and is closer to the continuum value  $-1$  than the  $6^4$  result. This is a finite-resolution regression check, not a claim of a converged integer at  $8^4$ . Independent Python/native comparisons additionally check gauge changes and reversal of parameter orientation.

For archived CP2K state exports covering skew cells, multiple centers, SOC spinors and reciprocal-boundary links, the largest difference between native and independently evaluated Gaussian overlaps is  $1.03 \times 10^{-15}$  with either one or two OpenMP threads. A cross-geometry overlap agrees to  $5.56 \times 10^{-16}$ . Table S2 reports physical atomic-translation pumps. The two-parameter case varies one reciprocal coordinate and one periodic atomic displacement. The separable four-parameter case uses the ordered coordinates  $(k_x, k_y, \phi_x, \phi_y)$ .

Table S2: Native physical Gaussian-snapshot tests. The minimum singular value is measured before polar unitarization. The phase column reports the largest diagnostic plaquette phase in radians. The  $4^4$  separable case tests the four-parameter physical workflow, not generic  $C_2$  convergence.

| Mesh           | Unrounded pairing     | Minimum singular value | Maximum phase |
|----------------|-----------------------|------------------------|---------------|
| $4 \times 4$   | $C_1 = -1.0000000000$ | 0.234278               | 0.402344      |
| $8 \times 8$   | $C_1 = -1.0000000000$ | 0.609933               | 0.099083      |
| $12 \times 12$ | $C_1 = -1.0000000000$ | 0.780474               | 0.044040      |
| $4^4$          | $C_2 = -1.0000000000$ | 0.233877               | 0.402347      |

The native  $4 \times 4$  to  $8 \times 8$  physical refinement check passed at an invariant tolerance of  $10^{-8}$ . The independent comparison-tool suite passed 34 tests without Z2Pack, with five optional-adaptor tests skipped, and all 39 tests with Z2Pack installed. The accompanying `validation/topology/native-phasons/` directory contains the machine-readable overlap and pump results, test logs, and source-version record. These are implementation tests and introduce no new material prediction or change to the neon/stanene results.

## S4 Native Inversion Representations and Symmetry Signatures

The inversion extension was validated separately from the phason extension and from the earlier Si/Al and Wilson calculations. Its source patch, base snapshot, inputs and outputs remain in `validation/topology/inversion-tqc/`. These records identify the original tests without retrospectively assigning them to another executable.

The WANNIER90 print section accepts KPOINTS\_SOURCE TRIM, SOC T, TIME\_REVERSAL T and INVERSION\_TQC T. TQC\_DIMENSION selects the four  $k_z = 0$  TRIM or all eight three-dimensional TRIM. PARITY\_ORIGIN specifies the inversion center in fractional direct coordinates. The occupied spinor bands must be a complete lowest-band prefix, with excluded states available for checking spectral separation. A PARITY-only mode also accepts isolated scalar subspaces. Neither route requires Z2Pack, Wannier90 execution or a Wannier fit. The TRIM calculation holds the converged SCF potential fixed and does not require the SCF integration mesh to contain those points. The two-dimensional option samples the fractional  $k_z = 0$  plane of a three-dimensionally periodic cell, not a separate two-dimensional Poisson boundary condition. The inversion center is specified explicitly rather than inferred from a space-group search.

For an orbital of angular momentum  $l$  on atom  $a$ , inversion about  $\mathbf{c}$  maps it to an identical Gaussian orbital on atom  $b$  when  $2\mathbf{c} - \mathbf{r}_a = \mathbf{r}_b + \mathbf{L}$  with integer  $\mathbf{L}$ . The coefficient multiplier is  $(-1)^l \exp(2\pi i \mathbf{k} \cdot \mathbf{L})$ . The code verifies atom-kind matching, involution, metric covariance,  $C^\dagger S C = I$ , unitarity and Hermiticity of the projected inversion matrix, and its commutator with the selected eigenvalues. Spinful time reversal is checked through the sewing matrix of  $\Theta = i\sigma_y K$ . Parity multiplicities must be even before division into Kramers pairs. The dimensionless and geometry tolerances are  $10^{-6}$  and  $10^{-6}$  bohr, respectively. The energy-commutator tolerance is

$10^{-6}$  hartree.

For  $N$  occupied Kramers pairs and odd counts  $n_{\kappa}^{-}$  at  $\mathbf{k} = \kappa/2$ ,  $\kappa \in \{0, 1\}^d$ , the parity index is  $\nu = \sum_{\kappa} n_{\kappa}^{-} \bmod 2$ . In three dimensions the weak indices restrict this sum to  $\kappa_i = 1$ , and the convention for the inversion indicator is  $z_4 = \sum_{\kappa} n_{\kappa}^{-} \bmod 4$ . The analytic inversion-subgroup EBRs are local even or odd Kramers pairs at centers  $\mathbf{c} + \mathbf{a}/2$ . Define

$$d_{\mathbf{a}} = 2^{-d} \sum_{\kappa} (-1)^{\mathbf{a} \cdot \kappa} (N - 2n_{\kappa}^{-}). \quad (\text{S4})$$

An integer atomic-signature decomposition exists when all  $d_{\mathbf{a}}$  are integers and  $N - \sum_{\mathbf{a}} d_{\mathbf{a}}$  is even. A nonnegative decomposition additionally requires  $N \geq \sum_{\mathbf{a}} |d_{\mathbf{a}}|$ . The implementation makes these decisions with integer arithmetic. Atomic compatibility concerns the symmetry signature, not a proof of trivial topology. Likewise, a signed-only decomposition does not by itself establish fragile topology. This inversion-only EBR test is distinct from the subsequent native little-group and explicit-segment analysis described below; it does not become a general EBR classifier when that additional analysis is used.

Table S3: Additional native parity checks with four occupied Kramers pairs. The four columns refer to  $(k_x, k_y) = (0, 0), (1/2, 0), (0, 1/2), (1/2, 1/2)$  at  $k_z = 0$ . Coordinates of the chosen inversion center are fractional.

| System  | $\mathbf{c}$      | $n_{\kappa}^{-}$ |   |   |   | $\nu$ | Atomic signature |
|---------|-------------------|------------------|---|---|---|-------|------------------|
| Neon    | $(1/2, 1/2, 1/2)$ | 3                | 3 | 3 | 3 | 0     | Compatible       |
| Stanene | $(0, 0, 0)$       | 1                | 2 | 2 | 2 | 1     | Obstructed       |

These tests retain the basis, cutoff, geometry and SCF meshes of the earlier material examples, replacing the Wilson paths by TRIM sampling. The three-dimensional neon test, with inversion center at the origin, also gives zero strong and weak indices and  $z_4 = 0$ . All seven inputs in the extended `QS/regtest-topology` set passed with two MPI ranks and two OpenMP threads per rank. The stanene parity result was independently reproduced with one rank and one thread. A single-rank, multiple-thread run exposed inconsistent state normalization in the existing local k-point diagonalization path for this build. The new metric checks reject that case. This extension does not resolve or validate that threaded path.

The lightweight CI fixture subsequently reduces the stanene SCF mesh from  $8 \times 8 \times 1$  to  $4 \times 4 \times 1$  and suppresses unused restart files. It still evaluates all four TRIM and reproduces the odd-pair counts  $(1, 2, 2, 2)$  and  $\nu = 1$ . This fixture is distinct from the archived  $8 \times 8 \times 1$  manuscript calculation and does not replace it or establish k-mesh convergence.

The native unit test covers multidimensional-character decomposition, nonorthogonal metrics, degenerate-state gauge changes, broken-symmetry rejection, all inversion-center EBRs and strong,  $z_4 = 2$  and signed-only examples. Independent enumeration of all two-Kramers-pair atomic sums matches the solver for all 81 two-dimensional and 6,561 three-dimensional candidate parity signatures. A deliberately incorrect inversion center is rejected in an end-to-end calculation. TRIM sampling does not prove a gap over the full Brillouin zone, and the SOC subspace still requires its usual convergence checks.

## S5 Native Little-Group and Compatibility Validation

The native algebra constructs unitary projective character tables, Wigner corepresentation restrictions, and endpoint-to-segment branching without a runtime dependency on Spgrep or a pretabulated irrep catalogue. The electronic-structure path uses SPGLIB to detect spatial operations and, when requested, adds their nonmagnetic time-reversal extension. It retains centering translations and tests the metric, band-subspace closure, energy commutators and the complete semilinear multiplication law before accepting band multiplicities.

The all-setting sweep covers all 530 Hall settings of the 230 ordinary space groups. At each setting, it samples  $(0, 0, 0)$ ,  $(1/2, 1/2, 1/2)$ ,  $(1/2, 0, 0)$ ,  $(1/3, 1/3, 0)$ ,  $(0, 0, 1/4)$  and  $(0.173, 0.217, 0.319)$  with scalar and spinful factors. The invariant cell metric is obtained by averaging over the supplied rotations. Both the ordinary group and its grey extension are checked. Generic-site orbits with physical spin multiplicity provide fixed-rank compatible band data. Segments run from Gamma to the other sampled points and from Gamma to its reciprocal image  $(0, 0, 1)$ . This tests directed-image monodromy as well as ordinary endpoint restrictions; it does not sample every possible line or material.

| Checked object                                        | Successful cases |
|-------------------------------------------------------|------------------|
| Ordinary/grey scalar/spinful factor systems           | 12,720           |
| Complete unitary projective character tables          | 6,360            |
| Endpoint tables with an antiunitary stabilizer        | 5,324            |
| Ordinary/grey directed segment comparisons            | 12,720           |
| Segments retaining an antiunitary stabilizer          | 4,462            |
| Gaussian directory assertions, serial / two MPI ranks | 69 / 69          |

Table S4: Native representation validation on 24 September 2026. The counts refer to algebraic cases or numerical assertions, not distinct material calculations.

Analytic cases separately test magnetic half translations, spin quaternion doublets, complex conjugate  $C_3$  pairs, twisted changes of factor gauge, and invalid pairing/doubling counts. Segment cases include displaced-origin PT pairs, endpoint-only time reversal, magnetic-glide Kramers planes, and spinful glide partner switching: a two-band hourglass connection is obstructed, whereas the complete four-band space is compatible. Unitary screw and mirror branching, reciprocal-image monodromy and changed coset representatives are also checked. The complete Hall-setting suite passes in release and with bounds checking and floating-point traps enabled.

The Gaussian directory uses UZH MOLOPT bases for its GPW fixtures and a contracted all-electron basis for GAPW. Ne tests exercise scalar and post-SCF GTH-SOC multiplets at Gamma, a generic PT-symmetric point and nonsymmorphic screw sectors. Scalar He/GPW, He/GAPW and Ne/SOC close a four-segment reciprocal loop; a separately selected single screw branch must report its monodromy obstruction. The generated Wilson-link test also exercises repeated refinement. Serial runs use two OpenMP threads; MPI runs use two ranks with two threads each and the thread-safe BLAS settings stated below. An earlier wider corepresentation integration pass additionally gave 110/110 serial and 111/111 MPI assertions across topology, Kubo, localizer and quadratic directories. These are separate test scopes, not numbers to add into a material benchmark total.

A separate local algebra sweep uses SPGLIB 2.7 magnetic data in the supplied setting for every one of the 1,651 magnetic UNI numbers. At the same six points and two spin choices, all 19,812 factor cases pass; 14,088 contain an antiunitary stabilizer, and 5,724 do not. For 10,154 primitive-cell tables supported by an independent Spgrep 0.7.0 comparison, dimensions, Wigner types and unitary restriction characters agree within  $7.11 \times 10^{-15}$ . Both calculations use the same supplied factor system, so this is not an independent validation of symmetry detection. Moreover, a separate semilinear-product check did not validate every explicit antiunitary matrix produced by that reference calculation. Those matrices are not imported into CP2K; the comparison is limited to types, dimensions and unitary restrictions. CP2K’s actual Gaussian sewing matrices pass their own full product check.

The separate native site-induction test samples fractional positions  $(0, 0, 0)$ ,  $(1/2, 0, 1/4)$  and  $(0.173, 0.217, 0.319)$  at the same six reciprocal points. Ordinary and grey groups, each with scalar and spinful factors, give 38,160 site/k-point cases over all 530 Hall settings. There are 67,824 local-representation columns evaluated and decomposed, and the same number of directed compatibility comparisons. Their atomic band rank equals the orbit multiplicity times the local irrep or corepresentation dimension. This construction preserves conventional-cell translation

sectors rather than silently converting them to a primitive-cell catalogue.

An independent check constructs Cartesian s/p/d orbital traces and, for spinors, multiplies them by the physical  $SU(2)$  trace. After matching the onsite orbital content, its explicitly translated orbit trace agrees with the corresponding induced character sum in 114,480 comparisons. The numerical acceptance tolerance is  $10^{-7}$  for these summed traces and decomposition checks; group and orbit geometry use  $10^{-9}$ . Analytic checks additionally cover inversion-center parities, screw orbits, spin-quaternion doublets, and complex local orbitals related by an antiunitary half translation. These are algorithmic reference tests, not new material calculations or a new production input mode. Additional scalar and spinful checks preserve the characters under passive origin shifts and unimodular changes to a sheared cell, and verify the expected Bloch rephasing under integer changes of operation representatives. The full database sweeps pass in serial, two-rank MPI, and a separate bounds-checked build with floating-point traps.

The native Wyckoff generator separately constructs 3,467 affine site-symmetry families across all 530 Hall settings and determines 5,648 directed specialization relations. External validation uses SPGLIB’s tabulated `Wyckoff.csv` coordinates, not the production fixed-set algorithm, to check dimensions, orbit multiplicities, family equivalence and closure containment. Every family matches bijectively, with largest affine residual  $2.22 \times 10^{-16}$ . The reference table’s version and checksum, comparison script and results are retained in the development record. No tabulated coordinates are read by the CP2K implementation.

For each generated seed the separate spinful induction kernel confirms its orbit and site stabilizer. Combined arbitrary origin shifts, nonorthogonal unimodular cell changes, integer shifts of operation representatives and reordered operations preserve the families and their specialization graph. Analytic tests cover inversion centers, intersecting mirrors, fixed-point-free screw axes, grey groups and an antiunitary half translation. Failure to resolve a generic stabilizer or an arithmetic/work limit returns no partial list. Maximal site groups are identified geometrically; this is not an elementarity test for their induced representations. The complete Wyckoff sweep passes in serial release, two-rank MPI and bounds-checked debug builds.

The subsequent atomic-signature module combines all generated sites and onsite irreps/corepresentations over the six reciprocal points listed above. Ordinary/grey and scalar/spinful variants yield 2,120 matrices, 31,079 provenance-preserving columns and 707,038 integer entries. Seven directed comparisons per column include the supplied point connections and a reciprocal-image loop, totalling 217,553 column/segment checks. Each block reproduces the column’s band rank; antiunitary rows count full corepresentations rather than treating all states as Kramers pairs. Additivity, permuted operation rows, and integer coset rephasing are checked. All sixteen inversion-center signatures and the generic-site signature are recovered analytically.

These full sweeps pass in serial release, two-rank MPI and debug with bounds and floating-point checks. A local optimized-compiler issue in automatic `MATMUL` result resizing was independently reproduced without CP2K and avoided by explicitly dimensioning the test vectors; no physical reference values were adjusted. The compiler reproducer and validation logs are retained separately from material results. No new general-classification input mode, elementarity assertion or complete reciprocal-stratum coverage follows from these matrix tests.

The integer-signature extension factors these same 2,120 matrices using checked 64-bit Smith reduction. Each factorization verifies its full integer certificate. One signed target per matrix is reconstructed, and the nonnegative search recovers every atomic column and a three-column sum per matrix: 33,199 witnesses in total. These deliberately constructed targets test recovery, not a classification of DFT bands.

An independent unit program compares Smith invariants with the gcds of all minors and computes the determinants of both unimodular transforms for 384 matrices with dimensions one through four. The rank-zero through rank-four case counts are 38, 167, 103, 60 and 16. The analytic inversion Fourier criterion supplies a separate exact oracle for every odd-parity-count pattern with two Kramers pairs: among 81 four-TRIM patterns, 41 admit signed and 33 admit nonnegative atomic signatures; among 6,561 eight-TRIM patterns, the counts are 241 and 129.

Thus signed-only cases are tested explicitly. For every  $2 \times 3$  matrix with entries in  $\{0, 1, 2\}$  and every two-component target in  $\{0, \dots, 4\}^2$ , a brute-force enumeration of coefficients zero through four supplies another oracle, covering 18,225 systems. Coefficients beyond four cannot contribute to a target in this range through any nonzero nonnegative column. Zero columns, deficient rank, invalid data, unsafe arithmetic and exhausted node budgets have distinct tests. The latter return an unresolved decision, not a symmetry obstruction.

The expanded matrix tests pass in optimized serial and two-rank MPI runs, and with the integer kernel compiled using bounds checks, floating-point traps and undefined-behavior instrumentation. MPI repeats the native algebra on each rank and does not distribute the catalogue construction. The analytic matrix path also passes with the test's SPGLIB preprocessor option disabled, linked against the existing library; this is not a full no-SPGLIB rebuild.

The Gaussian connection is enabled by `ATOMIC_SIGNATURES`. It compares only physically validated, sampled band subspaces with the generated reference matrix, preserving conventional-cell sectors. Band characters are rephased for integer changes of operation representatives, and the entire product/factor correspondence is checked, including antiunitary conjugation. All 2,120 algebra cases recover the same signature after operation and irrep permutations and such coset shifts. The shifted factors are regenerated by the separate little-group constructor rather than copied from the mapping formula. Inconsistent factors and unmatched operations are rejected. These expanded sweeps pass in optimized serial, two-rank MPI and with the new mapping kernel under bounds and undefined-behavior checks.

In the physical fixtures, scalar Ne at Gamma, Ne/SOC at Gamma and a generic point, centered/sheared He, and complete He/GPW, He/GAPW and Ne/SOC screw branches all have verified signed and nonnegative atomic signature witnesses. The independently obstructed single He screw branch has neither. The refined Wilson fixture repeats the comparison on both its four- and twelve-point sets. An orthorhombic He TRIM fixture sets `ATOMIC_SIGNATURE_MAX_NODES` to zero: signed membership is established, while nonnegative membership remains explicitly unresolved. This separates absence of a witness from proof of absence. The Gaussian-integration driver passed 73 assertions serially and with two MPI ranks, consisting of four native unit programs and 69 Gaussian-directory assertions. Existing physical reference values are unchanged; the new fixture uses the same UZH MOLOPT/GTH family as the other GPW tests. An independent reader then checks the exported records using arbitrary-size integer arithmetic: each of the serial and MPI runs contains 11 matrices and 19 signed/nonnegative witnesses satisfying  $Ax = b$  exactly, with consistent row/column band ranks. This verifies the stored certificates rather than independently classifying the negative cases. The reader and compressed output records accompany `validation/development/gaussian-signature-validation.md`.

The subsequent reciprocal-geometry test enumerates fixed families of  $(-1)^a W^{-T}$  for all 530 Hall settings and their grey extensions. Across these 1,060 geometric groups it retains 18,700 connected fixed-torus closures, 57,357 specialization incidences and 18,730 primitive integer tangent cycles. Full original cosets are restored when checking spinful stabilizers. At each setting, independent direct-transpose tests on reciprocal meshes with denominators 2, 3, 4 and 6 assign every point to exactly one matching family; this gives 333,900 point checks. Joint direct-cell shears, origin shifts, integer coset changes and operation permutations preserve the family and incidence correspondences. Analytic cases separately cover time reversal,  $PT$ , magnetic half translations, screw-invariant lines, centering and orthorhombic intersections.

For each ordinary/grey group, scalar and spinful atomic reference matrices are constructed at the automatically generated seeds and incidence images. Every atomic column obeys each directed restriction and primitive-cycle relation, giving 2,481,572 column/edge comparisons. This explicitly retains reciprocal shifts, nonsymmorphic phases and antiunitary restrictions. The complete sweep passes in serial and on both MPI ranks; the MPI run repeats the algebra, rather than distributing it. The new reciprocal module and its test also pass with bounds checks, floating-point traps and undefined-behavior instrumentation; their dependencies use the existing library. A database-free analytic test build passes against that library, not against a full no-

SPGLIB rebuild. The updated focused driver passes 74/74 serial and 74/74 MPI assertions: five native programs and the unchanged 69 Gaussian assertions. Logs and reproduction details are in `validation/development/reciprocal-strata-validation.md`. These are the geometric-stage results; fixed closures and their incidence/cycle graph are not a complete equivariant open-cell complex or a global connectivity classifier.

### S5.1 Automatic Gaussian symmetry sampling

The subsequent `KPOINTS_SOURCE SYMMETRY` integration generates and refines those paths directly at the frozen SCF Hamiltonian. Four Gaussian fixtures cover scalar He/GPW with UZH MOLOPT/GTH, all-electron He/GAPW with UZH TZVPP functions, Ne/GTH-SOC with its complete scalar AO space in the second variation, and a unimodularly sheared He cell with consistently transformed coordinates. They each yield 15 fixed families, 56 sampled points and 82 directed segments, all compatible. Each has signed and nonnegative atomic-signature witnesses. The focused driver now passes 106/106 assertions serially and with two MPI ranks: five native programs and 101 Gaussian assertions. The 32 added assertions do not alter the previous energy or symmetry references.

An independent file-level replay reconstructs NNKP inputs from the recorded fractional points and reciprocal-image edges. The rectangular NNKP neighbour format needs extra, nonduplicate padding connections; comparisons of segment data use the original directed edges only. Across the four cases, selected eigenvalues differ by at most  $2.1 \times 10^{-13}$  eV and band characters by at most  $2.2 \times 10^{-14}$ . Endpoint-to-line restrictions, the full atomic signature matrix and its integer witnesses agree exactly. The native path creates no `.mmn` file, whereas the replay exercises the existing overlap exporter. This comparison tests the new sampling integration against the established explicit-point path, not an independent electronic-structure implementation.

Increasing the number of points per original segment from three to five gives 138 points and 164 compatible segments for He/GPW. Removing time reversal retains valid unitary restrictions and reports zero antiunitary segments. Negative tests reject a too-small point budget, an invalid subdivision count and a conflicting Wilson-loop request. Selecting one of the two He screw branches fails the isolation check at a generated degeneracy instead of yielding an atomic conclusion. Analytic subdivision tests independently check lifted endpoints, uniform increments and total winding for positive, negative and multiple boundary crossings, including budget failure without mutation. The full 1,060-group geometry sweep is repeated with path subdivision. Records and the replay procedure are retained in `validation/development/automatic-sampling-validation.md`.

### S5.2 Joint compatibility lattice and independent certificates

The following baseline precedes the star-identification extension described in the next subsection; its archived certificates are retained. The next extension assembles all endpoint-to-line restrictions into one integer matrix  $C$ , expressed in the same irrep/corepresentation coordinates as the atomic matrix  $A$ . In particular, a reciprocal-image self-edge subtracts two restrictions in the same point block and retains nonsymmorphic monodromy. An exact Smith kernel  $K$  and integer atomic coordinates  $Y$  satisfy  $CK = 0$  and  $KY = A$ . The Smith form of  $Y$  supplies the free and finite factors of  $\ker_{\mathbb{Z}} C / \text{im}_{\mathbb{Z}} A$ .

The 530-setting ordinary/grey and scalar/spinful sweep verifies every generated atomic column in the joint kernel with zero quotient class. Small-group variants reverse all directed edges and reciprocal shifts, reject a deliberately incompatible atomic column and enforce a tiny memory budget without returning partial data. Separate integer tests include a nonsaturated constraint row, a free-plus-torsion quotient, zero-dimensional kernels and incompatible targets. The four/eight-TRIM inversion matrices with equal-rank constraints yield  $\mathbb{Z}_2$  and  $\mathbb{Z}_2^3 \times \mathbb{Z}_4$ , respectively. For all 81 and 6,561 two-pair parity patterns, zero class is equivalent to the independently tested signed atomic membership.

The full sweep passes in optimized serial, on both MPI ranks and with the new integer/compatibility and reciprocal kernels compiled with bounds, floating-point traps and undefined-behavior checks. All printed graph ranks/factors agree. The instrumented database-free analytic subset also passes against the existing library, not a full no-SPGLIB build. MPI repeats the algebra independently on each rank.

An independent SymPy 1.14.0 reader validates eleven exported Gaussian quotient records from each of the serial and two-rank MPI runs. It checks  $CK = 0$ ,  $KY = A$  and  $CA = 0$ , computes the rational rank of  $C$  and an independent integer Smith form of  $K$  to certify saturation, and factors  $Y$  independently. It also checks unimodularity of the exported quotient row map. Dividing its nonzero transformed rows by their divisors must give a primitive full-row-rank lattice; free rows must vanish. Finally it solves for each compatible target in  $K$  and recomputes the exported residue vector using arbitrary-size integers. Incompatible targets must not have a quotient-class record.

For the automatic He/GPW, He/GAPW and sheared-He fixtures,  $C$  has 109 retained equations and 91 coordinates. Its kernel has rank six, no free quotient factor and one  $\mathbb{Z}_2$  factor. The Ne/SOC fixture has 81 equations, 67 coordinates, kernel rank five and quotient  $\mathbb{Z}_2^2 \times \mathbb{Z}_4$ . All four physical targets have zero class. These are deliberately small atomic examples, not nontrivial topological-material discoveries. The explicit four-point complete screw fixtures have trivial quotients. The selected single screw branch violates compatibility and is assigned no class. By contrast, the sparse explicit He and refined Wilson graphs have free ranks four and 23. Those ranks expose missing constraints or identifications in the supplied graph and are not labelled physical topological invariants. Even a zero free rank does not establish completeness of omitted congruence constraints or BZ-wide isolation.

The official focused and exporter drivers together pass 165/165 assertions in each of SSMP and two-rank MPI. This comprises 129 native and little-group assertions and 36 existing export/k-point assertions; 23 new checks inspect compatibility, kernel rank, free/cyclic factors and actual target classes without changing existing physical reference values. Reproduction commands, the independent reader and compressed certificates accompany `validation/development/compatibility-lattice-validation.md`.

### S5.3 Star transport and published quotient comparison

The subsequent extension includes star-arm identifications in the same integer matrix, using the star-character transport formula derived in the main text. A threefold screw example checks genuinely complex characters, antiunitary conjugation, reciprocal shifts and independent integer changes of destination representatives. A star-only graph reduces nine local coordinates to three independent ones without requiring any line edges; a deliberately mismatched target is rejected. Corrupted factors, nonfinite characters and unsafe phases must not return a permutation. Every retained star link in the complete Hall-setting sweep agrees with independently induced atomic columns; inverse maps and all alternative connecting operations give the same full-corepresentation identification.

An independent development reader compares the resulting ranks and ordered nontrivial Smith factors with Tables I–IV and IX–XII of [Po, Vishwanath and Watanabe \(2017\)](#). It validates all 230 entries in each rank table and uses the SPGLIB Hall-to-space-group metadata solely to associate settings with the reference rows. All 2,120 ordinary/grey and scalar/spinful Hall-setting combinations agree, with no free factors on these automatically generated graphs. These are representation group comparisons, not 2,120 DFT calculations. They do not independently prove completeness of a user-supplied graph or the elementary nature of every generated atomic representation.

The Gaussian certificate reader also checks each exported star’s reciprocal geometry. The explicit He/GPW and He/GAPW screw examples now contain ten equations for nine coordinates, still with compatible rank three and trivial quotient. The Ne/SOC screw has five equations, five

coordinates and rank two. The refined He Wilson graph contains 76 equations and 75 coordinates: its kernel rank falls from 42 to 23 and its free quotient rank from 23 to four. The partial two-point He graph still has four free factors; the single-branch screw obstruction remains incompatible. No energy, character or physical reference value is adjusted to obtain these results. The automatic Gaussian fixtures retain the ranks and finite factors of the baseline.

Reproduction scripts, reference-comparison records and the new certificates are archived separately in `validation/development/star-transport-validation.md`. The star-transport stage passed 170/170 assertions in each of SSMP and two-rank MPI: 134 native/little-group checks and 36 existing export/k-point checks. Five new Gaussian assertions count star links, including the zero-link monodromy control and the six links of the refined Wilson graph. The full star-aware sweep output is identical between optimized SSMP, each MPI rank and the instrumented build. The database-free analytic variant passes as well, using the existing library rather than a full no-SPGLIB build.

These algebra sweeps do not implement magnetic-space-group detection, a magnetic non-collinear ground state, canonical irrep labels, a complete Brillouin-zone connectivity graph or a full elementary-band catalogue. Explicit-segment compatibility is necessary, not sufficient, for a gapped band subspace. The retained summaries and test logs are in `validation/development/`; no production material classification is inferred solely from a successful database sweep.

#### S5.4 Real-space induction and composite witnesses

The subsequent site-specialization extension evaluates finite site-group induction at lifted endpoints of the generated Wyckoff families. All 22,592 relations in the ordinary/grey scalar/spinful 530-setting sweep pass their independently generated band-character identities at eight reciprocal points, giving 180,736 checks. Transitive singleton equivalences and direct-sum witnesses also preserve every tested character and physical band dimension. The optimized SSMP, both MPI ranks and bounds/undefined-behavior instrumented executable produce identical 4,774-line result records. Intermediate-position tests cover inversion and mirror families, cell shear, shifted origin and integer coset changes. A separate analytic build omits the database interface; this uses the existing CP2K library rather than constituting a complete no-SPGLIB build.

The detected exceptional-space-group coverage matches Tables II–IV of [Cano et al. \(2018\)](#) in all 2,120 variants. A second comparison uses the independently archived Bilbao EBR data in `Crystalline.jl`. Every per-setting generator count agrees; counting each space group once gives 3,383 scalar and 2,258 spinful generators without time reversal, and 3,141 scalar and 1,616 spinful generators with time reversal. The archived dataset’s 2,258 differs from the aggregate 2,263 printed in the 2018 article. The comparison retains the dataset’s values and does not adjust a native count to obtain agreement. Neither these totals nor agreement of exceptional-space-group coverage establishes a canonical irrep-by-irrep identification.

For each of 15 Gaussian detail files per build, an independent arbitrary-precision reader verifies  $AE = A$ , the conserved band dimensions, the lifted endpoint geometry, every source stabilizer’s action on the entire displacement, singleton equivalence classes and recursive decomposition provenance. The physical fixtures include GPW, all-electron GAPW, post-SCF GTH SOC, sheared cells, screw sectors, TRIM and refined Wilson sampling. The official serial and two-rank MPI drivers each pass 171/171 assertions: six native unit programs, 129 Gaussian little-group assertions and 36 existing export/k-point assertions. The bounded new unit includes the spinless time-reversal exception at an  $S_4$  site in space group 84.

The reproducible comparison scripts, versioned reference hashes, native logs and Gaussian certificates are retained in `validation/development/site-induction-validation.md`. The reference tables are test oracles, not runtime dependencies or imported CP2K classification tables. These tests prove the returned decompositions and their integration, not a globally gapped material classification, canonical EBR labels or completeness of arbitrary reciprocal sampling.

## S5.5 Affine-reference and site-resolved EBR checks

The next reference layer uses the separately archived space-group generators and parametrized Wyckoff coordinates of Crystalline.jl. The test fixture closes the generators with exact rational arithmetic; the native reference path does not obtain operations from SPGLIB. All 1,731 named affine families match the generated closures, including their whole tangent spaces, dimensions and conventional-cell multiplicities. A second pass simultaneously changes the basis, origin, integer coset representatives and operation order and preserves every named specialization edge. The four ordinary/grey scalar/spinful variants give 6,924 affine correspondences.

After the conventional-cell centering factor is removed, the band dimensions of the 10,398 unreduced maximal-site columns agree with the independent EBR multisets separately at each Wyckoff position. No reference column is selected or discarded using the legacy `Decomposable` field. This test detects incorrect distributions among sites or band dimensions that total counts would miss, but does not assign canonical labels to equal-dimensional irreps. Optimized serial and two-rank MPI output is compared with a separately instrumented reference executable compiled without its SPGLIB database interface. The latter still links the existing CP2K library, and is not a complete no-SPGLIB CP2K build. Numerical guards additionally reject wrong tangents, repeated reference entries, incorrect multiplicities and nonfinite or unbounded coordinates. Reproduction scripts, source hashes and logs are retained in `validation/development/affine-reference-validation.md`.

## S5.6 Complete character correspondence with reference tables

The `irreptables` 3.1.0 package supplies independent single- and double-valued characters at the tabulated reciprocal points of all 230 space groups. The fixture reader uses the data files only, not IrRep’s character-generation implementation. It records a SHA-256 hash for each of the 460 source files. Centering is removed by an explicit primitive basis  $P$ , with  $W' = P^{-1}WP$ ,  $\tau' = P^{-1}\tau$  and  $k_{\text{native}} = -P^T k_{\text{table}}$ . The sign converts the table’s translation-character convention to the native  $\exp(-2\pi i k \cdot T)$  convention; the original character columns and spin lifts are retained. Without this conversion, central-projector audits reject 25 boundary-point tables. With the declared conversion, all 2,700 tables pass the independent audit. This is a convention check, not evidence of an error in IrRep.

Crystallographic coordinates printed to five decimal places are recovered as rational values with denominators at most 48 and changes strictly below  $5.1 \times 10^{-6}$ . Characters are not rationalized or rounded. A Cartesian intertwiner aligns the reference spin matrices with the integer rotations; unitary polar factors are used only to identify this coordinate frame, while comparisons retain the raw reference matrices. Native spin representatives are returned by the same factor-system construction that generated their signs. An independent recomputation after a different floating-point cell normalization is not used to infer those signs.

Every character column is compared, including distinct irreps with the same dimension. All 8,907 columns match; the largest residual is  $1.20921 \times 10^{-5}$ . The native construction tolerance remains  $10^{-9}$ ; the external comparison uses  $2 \times 10^{-4}$  to accommodate the printed reference precision. The 11,607 normalized case and label records are identical in optimized SSMP, separately on both MPI processes, and in a focused GNU Fortran build with bounds checking, floating-point traps and undefined-behavior instrumentation. The instrumented executable links the existing CP2K library; this is not a fully independent no-SPGLIB build or a multi-node scaling test.

The bounded native test covers ordinary and projective  $C_3$ , complex rephasings, label-permuting one-dimensional twists, interleaved antiunitary operations, and paired/doubled Wigner restrictions. An analytic body-centered three-screw boundary case also changes the Cartesian orientation, operation order and cell scale from  $10^{-100}$  to  $10^{100}$ . Inconsistent factor-gauge maps, repeated character columns, NaNs, unbounded values, and orthogonal dimension-preserving mixtures that are not representation characters are rejected. Full onsite EBR naming, arbitrary-cell standardization for named labels, and global connectivity are not implied by these reference correspondences. The

official SSMP and two-rank MPI drivers each pass 173/173 assertions: eight native unit programs, 129 Gaussian little-group checks and 36 existing topology/export/k-point assertions. No physical reference value was changed for this work. Scripts, provenance and numerical logs are retained in `validation/development/character-reference-validation.md`.

### S5.7 Onsite axes, spin phases and reference correspondence

The onsite matcher retains the seed stabilizer’s Cartesian rotations, antiunitary flags, product table and actual spin lifts. An explicitly supplied orthogonal map aligns the reference axes with the native axes. With its spin representative  $V$ , a unitary reference lift transforms as  $VUV^\dagger$ , whereas the matrix part of an antiunitary lift transforms as  $VAV^T$ . Here  $A$  includes  $i\sigma_y$  for physical time reversal. Comparing these matrices with the retained native lifts determines a single phase per operation. The product and factor tables and all character columns must then agree after this declared operation/phase map; phases are not fitted independently to different irreps. The interface returns both the column correspondence and its operation/phase provenance.

For an independent point-group test, the 460 scalar/spinful Gamma tables of `irreptables` 3.1.0 are reused, with their space-group translations removed. Thus the group fixes the chosen site: this is not an enumeration of every actual Wyckoff-site stabilizer in the 230 space groups. Three variants combine cell scales  $10^{-100}$ , 1 and  $10^{100}$ , basis shear, shifted origin, reversed operation order, integer coset changes and Cartesian rotation. Table S5 summarizes the resulting comparison. The  $10^{-9}$  native tolerance is unchanged; the external tolerance is  $2 \times 10^{-4}$  for the finite-precision reference rotations, spin matrices and characters.

| Onsite correspondence check                           | Result                 |
|-------------------------------------------------------|------------------------|
| Gamma reference tables / coordinate variants          | 460/3                  |
| Complete-table matches / matched column records       | 1,380/6,384            |
| Reference-normalized character entries per executable | 88,602                 |
| Largest combined external comparison residual         | $3.14 \times 10^{-5}$  |
| Largest cross-build character difference              | $8.44 \times 10^{-15}$ |
| Largest cross-coordinate character difference         | $9.77 \times 10^{-15}$ |

Table S5: Onsite point-group correspondence validated on 24–25 September 2026. Each MPI process repeats the complete audit; these counts do not measure distributed speedup or material coverage.

In several cubic spinful cases, 52 native-column/label records differ between release and instrumented builds. The corresponding complete characters, indexed by the returned reference labels and operations and divided by the explicit lift phases, agree to the tolerance shown above. Raw internal irrep indices therefore must not serve as cross-build physical identifiers. Both MPI processes retain the release ordering, but correctness does not depend on that coincidence.

The bounded analytic tests distinguish the three complex  $C_3$  characters and their scalar/spinful grey restrictions, test improper axis changes and complex spin-lift rephasings, and reject inconsistent characters, missing antiunitary spin factors, duplicate operations and nonfinite or unbounded input. A nonsymmorphic  $P6_3$  example has two sites and a  $C_3$  onsite subgroup. At  $k_z = 0, 1/4, 1/2$ , operations exchanging sites have zero induced trace; the remaining characters agree with the analytic onsite sum and screw-translation phase. This holds for scalar/spinful ordinary and grey variants after origin, basis, axes and operation-order changes. Separate magnetic tests distinguish an onsite  $C_2\mathcal{T}$  operation, whose spinful square is  $+1$ , from an antiunitary half translation exchanging sites.

The focused instrumented executable checks the symmetry, little-group, corepresentation, character-matching and site-induction modules with bounds checks, floating-point traps and undefined-behavior instrumentation. It omits its SPGLIB database test interface but links the existing release CP2K library; it is neither a fully instrumented Gaussian calcu-

lation nor a complete no-SPGLIB build. The official serial and two-rank MPI directories each pass 179/179 assertions after this extension, without changes to physical reference values. Scripts, input hashes, complete character records and driver logs are retained with `validation/development/onsite-reference-validation.md`. This establishes a convention-aware onsite matching kernel, not automatic standard labels in the Gaussian output, complete actual-site EBR identities or a material topology classification.

## S5.8 Actual Wyckoff sites and independent Bloch induction

The subsequent comparison retains the real fractional space-group translations. The independently archived affine families supply one generic point per family, formed from their origin and tangent parameters  $\sqrt{2}/7$ ,  $\sqrt{3}/7$  and  $\sqrt{5}/7$ . An explicit integer-Hermite construction removes centering and records the primitive basis. All 1,731 conventional families are retained, including nonmaximal sites, rather than selecting a subset by EBR dimension or by the archive’s decomposability field.

The reference reader uses Gamma character tables for 32 symmorphic point-group representatives, separately for scalar and spinful representations. It matches their Cartesian operation sets to each actual site stabilizer by an explicit orthogonal conjugacy and records the chosen axes. For grey groups, independent character conjugation and the spin-aware Frobenius–Schur indicator construct paired or doubled unitary restrictions. These reference axes are not automatically the standard Bilbao onsite-label convention. The native matcher must accept the complete table, including equal-dimensional columns, in both the original and simultaneously rotated, sheared, origin-shifted, reordered and integer-coset-shifted coordinates.

An independent Python calculation then builds the real-space orbit and its representatives. For each unitary operation that fixes an orbit site modulo a translation, it evaluates the onsite character with an explicit physical  $SU(2)$  intertwiner and multiplies it by  $\exp(-2\pi i k \cdot L)$ . Summing these contributions gives the whole-band trace without using CP2K’s cached induction coefficients. The primitive reciprocal points are  $(0, 0, 0)$ ,  $(1/2, 0, 0)$ ,  $(0, 1/2, 0)$ ,  $(0, 0, 1/2)$ ,  $(1/4, 1/4, 1/4)$  and  $(0.173, 0.217, 0.319)$ . Only unitary little-group operations are compared at each point; their complete membership is checked independently. The native results are converted using explicit spin-lift and integer-coset phases, not by fitting a phase to each band representation.

| Actual-site and induction comparison                       | Result                  |
|------------------------------------------------------------|-------------------------|
| Wyckoff families / symmetry variants / coordinate variants | 1,731/4/2               |
| Complete onsite-table matches                              | 13,848                  |
| Onsite columns before coordinate repetition                | 16,957                  |
| Induced tables / character entries per executable          | 83,088/1,384,328        |
| Largest absolute induced reference difference              | $1.4510 \times 10^{-4}$ |
| Largest dimension-scaled reference difference              | $1.05 \times 10^{-5}$   |
| Largest induced cross-build difference                     | $1.87 \times 10^{-14}$  |
| Largest induced cross-coordinate difference                | $2.67 \times 10^{-14}$  |

Table S6: Complete actual-site comparison on 25 September 2026. The external differences reflect the finite-precision reference characters and spin matrices. Native construction uses  $10^{-9}$ ; the external bound is  $2 \times 10^{-4}$  after division by each induced column’s positive band dimension.

The optimized SSMP, both MPI processes and the focused instrumented executable pass all cases. The 189,664 onsite character entries agree across builds within  $8.22 \times 10^{-15}$ ; the larger induced comparison is summarized in Table S6. A deliberate phase reversal at site  $c$  of space group 100 changes one scalar induced character from  $-2i$  to  $+2i$  at the quarter-diagonal point. The test rejects it with absolute error 4 and dimension-scaled error 1. A separate output audit caught and corrected list-directed Fortran parsing of labels such as `2*GM1`; these labels are now read literally, without changing numerical reference characters.

This optional reference sweep adds no external database dependency or input keyword

to production calculations. The standard native unit tests also pass in serial and MPI with two OpenMP threads; the preceding 179/179 Gaussian/export driver results are not counted as new DFT runs for this test-only extension. The focus remains generic representatives and six sampled reciprocal points, not every coordinate in a free family, complete reciprocal-stratum connectivity, magnetic-space-group site enumeration or automatic canonical EBR naming. The archive, source/fixture hashes and replay scripts are recorded in `validation/development/wyckoff-character-validation.md`.

## S5.9 Recoverable onsite provenance in Gaussian output

The atomic-signature output retains the local character columns, their full site-group operation mapping, Cartesian axes, atom-cell shifts and physical spin lifts, together with the lifts of the complete space group. These data are copied from the construction used for the induced columns. No second choice of quaternion signs is made. The onsite product table is written as signed operation indices: its factors are exactly  $\pm 1$  at Gamma, which the writer checks before encoding. This reduces output volume without rounding characters or suppressing validation. Scalar characters do not use the recorded spin-1/2 matrices.

An independent reader reconstructs the orbit sums using only these records, then compares them with the little-group decompositions of the exported atomic columns. For each of the optimized serial and two-rank MPI builds, 16 retained Gaussian test outputs cover 197 site tables, 697 onsite columns and 263 reciprocal queries, giving 31,537 induced character entries. GPW, GAPW, post-SCF SOC, sheared cells, screw phases and automatic reciprocal sampling are included. The largest reconstruction error is  $3.58 \times 10^{-14}$ . All onsite columns also match the external point-group references, with largest combined lift/character residual  $9.56 \times 10^{-6}$ ; the declared axes and column correspondences agree between builds. These are reference-frame identifications, not canonical Bilbao onsite or EBR labels. Deliberate changes to a character, lattice shift, full-group spin entry or product sign are rejected in all 128 negative checks. Both regression drivers pass 179/179 checks; the high-cutoff Ne case uses a 64 MiB OpenMP thread stack to avoid a separately identified local grid-kernel stack overflow. No physical reference value or tolerance is changed. The native snapshot audit additionally passes all 2,120 Hall-setting/symmetry variants. The evidence and replay tools are retained in `validation/development/onsite-export-validation.md`.

## S5.10 Site-resolved named EBR character sequences

The EBR vectors supplied with `irreptables` 3.1.0 are combined with its independently tabulated little-group characters, without using the native induction algorithm to create the reference. The ordinary-group scalar and spinful datasets contain 1,599 site/spin cases and 5,641 named reference columns across all 230 space groups. There are 9,667 reciprocal queries and 297,529 induced character entries per coordinate convention. Every reference column must match the same native onsite column at all supplied points and operations. Candidate sets are retained, and their multiplicities are checked so that duplicate reference columns cannot consume one native column twice.

The earlier little-group fixture lacks some opposite-valley points. Explicit type-I tables supply 62 such points in 40 scalar/spinful tables, including KA/HA and PA/WA. Before use, the adapter verifies the complete operation map, physical spin lifts and integer coset shifts; all shared labelled characters agree within  $3.01 \times 10^{-16}$ . Missing labels are not assigned by guessed complex conjugation. The reference uses the already validated  $\mathbf{k}_{\text{native}} = -P^T \mathbf{k}_{\text{table}}$  convention and retains the printed character precision.

Optimized serial, both processes of a two-rank MPI run and a focused instrumented executable pass the original and combined sheared/rotated/origin-shifted/reordered conventions. They agree on all candidate multiplicities. Within the specified Wyckoff position, 5,263 reference columns have one candidate and 378 have multiple candidates. The largest dimension-scaled character residual is  $1.05 \times 10^{-5}$ , below the unchanged  $2 \times 10^{-4}$  reference tolerance; native construction still

uses  $10^{-9}$ . The MPI processes repeat the audit independently, so these results are not a scaling measurement.

Six deliberately malformed P4 spinful references are rejected by both optimized executables: a boundary character, a shifted site with unchanged Gamma data, a duplicated character column, a zeroed spin lift, a duplicate operation, and unexpected trailing data. The unchanged baseline passes. Instrumentation covers the test and the symmetry, little-group, corepresentation, character-matching and induction kernels; other modules remain release-library code. This is not a fully instrumented Gaussian or MPI calculation. The generated local space contains 5,704 columns at the selected sites; 63 are absent from the supplied EBR references. Their elementarity, grey-group reference matching, canonical onsite axes and complete band connectivity are not certified by this test. Scripts, input and source hashes, full output and replay instructions are retained in `validation/development/ebr-character-validation.md`.

### S5.11 Character functions beyond finite reciprocal samples

Equal characters at selected reciprocal points need not imply equal character functions. We distinguish these cases without replacing the sampled integer signature matrix or its real-space induction witnesses. For a unitary operation with integer direct-space rotation  $W$ , the invariant reciprocal torus is

$$K_W = \{\mathbf{k} \bmod \mathbb{Z}^3 : (W^T - I)\mathbf{k} \in \mathbb{Z}^3\}. \quad (\text{S5})$$

Its character lattice is  $\mathbb{Z}^3/(W - I)\mathbb{Z}^3$ . Indeed, two integer shifts differing by  $(W - I)\mathbf{n}$  give identical Bloch phases on  $K_W$ . For the checked Smith factorization  $U(W - I)V = D$ , each finite induction term  $c_j \exp(-2\pi i \mathbf{k} \cdot \mathbf{L}_j)$  is therefore collected under the key  $U\mathbf{L}_j$ , modulo each nonzero Smith divisor. Rows with zero divisor retain an unrestricted integer key. The coefficients  $c_j$  include the local character, the spin-lift factor and, where required, conjugation by an antiunitary orbit representative. Antiunitary operators themselves have no basis-independent trace and are not assigned character functions.

These quotient characters are linearly independent on the complete compact group  $K_W$ . Equality of the collected coefficients thus tests equality of the finite character functions, including all disconnected torsion components, rather than fitting more  $\mathbf{k}$  points. Integer operations are checked for overflow; complex coefficients are compared at  $10^{-9}$  after scaling by band dimension in this audit. The production diagnostic uses the supplied character tolerance and rejects nontransitive tolerance neighborhoods. It reports local unitary-character classes but does not merge atomic columns, change membership decisions or infer an intertwiner between Bloch bundles.

As an analytic counterexample to Gamma-only matching, equal-parity orbitals at inversion centers  $(0, 0, 0)$  and  $(1/2, 0, 0)$  have the same Gamma characters but different characters on other components of the inversion-fixed torus. The Fourier comparison distinguishes both parity pairs. With only Gamma supplied to the atomic catalogue, all sixteen centered inversion columns and the generic-orbit column remain distinct, also with Kramers pairing.

For the 378 multi-candidate references in the preceding 230-space-group EBR audit, however, every candidate set has equal collected unitary-character coefficients in both tested coordinate conventions. Optimized serial and both MPI processes agree. These ambiguities are therefore not removed by denser sampling; canonical onsite conventions and bundle-level information remain separate requirements. Independent evaluation of the collected Fourier sums against raw Bloch induction checks 213,238 operation/ reciprocal-point queries per optimized executable or MPI process, with a largest dimension-scaled discrepancy below  $4.1 \times 10^{-15}$ . This count includes the bounded native database and analytic controls, not additional material calculations or an MPI scaling measurement. Full logs and replay instructions are retained with the evidence below. A separate all-setting sweep checks 536,410 Fourier evaluations for all 530 Hall settings, ordinary/grey and scalar/spinful groups at three sites, with a largest discrepancy below  $3.8 \times 10^{-16}$ . The updated atomic catalogue constructs 2,120 matrices with 31,079 columns in both optimized builds, while retaining its signed and nonnegative induction witnesses. The 179 relevant native, Gaussian, topology

and k-point regression checks pass in each build. Independent reconstruction from 16 Gaussian detail files per build retains the earlier 31,537 character-entry comparison and rejects 128 deliberately corrupted records across both builds. The focused instrumented run additionally checks 205,614 Fourier evaluations; it includes the induction and checked integer kernels but not the native SPGLIB database sweep or a fully instrumented Gaussian calculation. Data, source hashes and replay instructions are retained in `validation/development/fourier-character-validation.md`.

The Gaussian exports also illustrate why the two diagnostics must remain separate. The counts below refer to generated atomic reference columns, not numbers of occupied bands or Wannier functions. Both optimized builds give the same counts; the largest accepted coefficient residual in these exports is  $1.86 \times 10^{-17}$ .

| Test input                      | Columns | Sampled classes | Full-torus classes |
|---------------------------------|---------|-----------------|--------------------|
| <code>helium.inp</code> (GPW)   | 125     | 18              | 125                |
| <code>helium-gapw.inp</code>    | 11      | 5               | 11                 |
| <code>neon-soc.inp</code>       | 33      | 17              | 29                 |
| <code>neon-spin-lift.inp</code> | 26      | 23              | 23                 |

Here full-torus classes still mean equality of unitary characters only. Their identifiers are local to each constructed catalogue; all original columns and the sampled membership results are retained.

### S5.12 Shared spin lifts in the Gaussian SOC analysis

The Gaussian sewing matrices must use the same  $SU(2)$  representatives as their factor table. This is tested with Ne in the primitive body-centered cubic basis

$$A = 4.2 \begin{pmatrix} 1 & 0 & 1/2 \\ 0 & 1 & 1/2 \\ 0 & 0 & 1/2 \end{pmatrix} \text{ \AA}, \quad (\text{S6})$$

with roundoff-sized perturbations of the input components. The calculation uses PBE, the DZVP-MOLOPT-GGA-GTH-q8 basis from `BASIS_MOLOPT_UZH`, post-SCF GTH SOC, all 13 scalar states in the second variation, and the lowest isolated Kramers doublet as the represented subspace. Both the SCF and the property calculation use Gamma. A single 1200 Ry real-space grid resolves the quadrature sufficiently for the unchanged  $10^{-6}$  metric/subspace and  $10^{-6}$  hartree energy-commutator tolerances.

For identical inputs, the independent-lift and shared-lift paths have metric residual  $1.21 \times 10^{-15}$  and subspace residual  $2.90 \times 10^{-7}$ . Their product-law residuals are respectively 2 and  $8.49 \times 10^{-14}$ . The former correctly aborts; the latter identifies one two-dimensional unitary irrep and one undoubled corepresentation, with signed and nonnegative atomic membership. The SCF energy is  $-34.905537447716$  hartree in both runs, differing by less than  $3 \times 10^{-14}$  hartree before rounding. There is no post-hoc fitting of sewing matrices, reference-energy change, or relaxation of the symmetry checks.

Two distinct convergence issues were separated from the lift defect. At 400 Ry, the Gamma calculation has subspace residual  $9.19 \times 10^{-5}$  and fails the closure test before reaching the product test. Also, the shifted  $2 \times 2 \times 2$  Monkhorst–Pack set in this particular primitive basis is preserved by only 8 of the 48 cubic spatial operations; increasing its real-space cutoff does not restore the missing integration-set symmetry. These diagnostic failures are retained, not converted to successful references. The shared-lift native test additionally applies all 48 cubic operations and their 48 time-reversed partners to complex frames, including cell scales from  $10^{-100}$  to  $10^{100}$ . With this additional Gaussian test the official SSMP and two-rank MPI drivers each pass 179/179 assertions. The focused spin/factor unit tests also pass with bounds checking, floating-point traps and undefined-behavior instrumentation; this is not a fully instrumented Gaussian or MPI build. The reproduction record is `validation/development/gaussian-spin-lift-validation.md`.

## S6 Bloch Kubo and Property-Symmetry Validation

The retained development report is `validation/development/bloch-response-summary.md`. It records distinct finite-cutoff regression fixtures and the higher-cutoff primitive/supercell comparison. For Ne with post-SCF GTH SOC and PBE at 2000 Ry, the  $3 \times 1 \times 2$  primitive mesh gives  $-34.923548219306689$  hartree, compared with  $-209.541289276241002$  hartree for six cells sampled explicitly at Gamma. Both print  $6.21105320 \times 10^{-9}$  S. At the inexpensive 200 Ry PADE regression settings, the respective conductivities are  $6.42389197 \times 10^{-9}$  and  $6.35401974 \times 10^{-9}$  S. Those distinct references are retained: the mismatch is not removed by changing a reference value.

Ne/SOC on a  $2^3$  property mesh uses four representatives. Bi<sub>2</sub>/SOC on  $3^3$  uses six; its printed nine-component tensor differs from full-grid by at most  $2.71 \times 10^{-15}$  relative to its largest entry. Primitive diamond/SOC on  $2^3$  uses three representatives with either K290 or SPGLIB. Its tensor difference is  $2.21 \times 10^{-11}$  in the same normalization. Metric, Hamiltonian and current reconstruction residuals are approximately  $2.1 \times 10^{-13}$ ,  $6 \times 10^{-11}$  and  $1.3 \times 10^{-11}$ , respectively. These are operator-covariance fixtures, not converged material transport predictions.

The final recorded directory drivers passed 132 serial and 162 MPI assertions. The MPI layout was four ranks with two OpenMP threads, and the serial build also used two OpenMP threads. Coverage includes moments, k-point/Wannier export, topology, Kubo, localizer, and quadratic directories; the MPI build additionally tests MUMPS and Tacho. Mathematical unit tests cover the constituent operator identities separately.

The Apple ARM runtime required one BLAS thread, a consistent thread-safe OpenBLAS preload, and an OpenMP stack of 128 MiB. Without these settings, failures occurred before the property code. An implicit-Gamma high-cutoff conversion also exposed a slow Miller-index lookup; the completed comparison uses an explicit Gamma input. These runtime qualifications must accompany the results, rather than be presented as production scalability.

## S7 Localizer and Quadratic Validation Records

The transferred developer notes in `validation/development/` document finite and periodic formulations, dense and sparse factorization, and their limitations. They are implementation records, not independent material benchmarks. The matrix tests include nontrivial and trivial Chern models, spin-mixed Kramers-pair models, origin and energy shifts, nonorthogonal congruences, invalid metrics, and gap closings.

Dense LDL tests cover both  $1 \times 1$  and  $2 \times 2$  blocks. Pfaffian tests independently cover odd permutations and recursive small-matrix reference polynomials. Sparse gap brackets are compared with generalized dense spectra, not with pivot sizes. Realification duplicates the spectrum and is accounted for in every inertia comparison.

The quantitative records below are retained calculations from 24 September 2026, not new material runs performed during manuscript editing. Their raw logs, molecular inputs, state coefficients and model-test source snapshot are archived in `validation/development/solver-test-records.tar.gz`. The accompanying `summarize_solver_tests.py` independently re-parses these records, checks the reported gap brackets and recomputes the molecular metric-subspace comparisons. It writes `solver-test-summary.json`, including SHA-256 checksums for all archived members. Model indices are asserted by the original unit driver; its performance lines do not print a separate index.

### S7.1 Sparse inertia: accuracy and indicative cost

The open square Qi–Wu–Zhang fixture has two orbitals per site, onsite block  $-\sigma_z$ , and positive-direction hopping blocks  $(\sigma_z - i\sigma_x)/2$  and  $(\sigma_z - i\sigma_y)/2$ . The query is at the lattice center and zero energy, with  $\kappa = 0.25$  in model energy per lattice spacing. All nine configurations in Table S7 give index  $-1$ . The dense generalized gap lies within the sparse bracket to the driver’s  $10^{-9}$

comparison tolerance, both before and after the tested nonorthogonal congruence. The printed gap is 0.81017072 for side length 10 and 0.81017480 for lengths 20 and 30. These are model energy units, not first-principles hartree energies.

| Side | Real order | Ranks | Dense (s) | MUMPS (s) | Fill ratio | Memory (MB) |
|------|------------|-------|-----------|-----------|------------|-------------|
| 10   | 800        | 1     | 0.03321   | 0.05853   | 15.0       | 1           |
| 10   | 800        | 2     | 0.03326   | 0.05922   | 15.2       | 2           |
| 10   | 800        | 4     | 0.03427   | 0.05682   | 15.1       | 4           |
| 20   | 3200       | 1     | 1.89561   | 0.51263   | 21.9       | 6           |
| 20   | 3200       | 2     | 1.88879   | 0.48799   | 22.3       | 9           |
| 20   | 3200       | 4     | 1.89937   | 0.33720   | 22.5       | 17          |
| 30   | 7200       | 1     | 22.94493  | 2.40084   | 36.0       | 19          |
| 30   | 7200       | 2     | 23.11979  | 2.20658   | 37.2       | 32          |
| 30   | 7200       | 4     | 23.23512  | 1.36989   | 36.4       | 45          |

Table S7: Single final runs of the open-model localizer kernel. Real order is twice the complex localizer order. Dense means source-rank generalized diagonalization plus the LDL cross-check; MUMPS includes symbolic analysis, metric checks and all gap shifts. Fixture and coordinate-list construction are excluded. Fill ratio is the maximum factor-entry count divided by input entries; memory is MUMPS’s process-summed solver allocation, not application RSS.

The recorded host is an Apple M4 Max with 14 cores and 36 GiB, using GNU Fortran 16.1, OpenMPI 5.0.9, OpenBLAS 0.3.33 and MUMPS 5.9.1 with PORD ordering. Both BLAS and OpenMP use one thread in these kernel measurements. The larger four-rank run recovers from insufficient factor workspace by increasing the workspace estimate, without changing matrix entries or pivot tolerances. The increasing fill ratio and summed memory make clear that sparse does not imply linear storage or ideal MPI speedup. The test retains dense oracle matrices, so its process RSS would not measure the production sparse path alone. No error bars or statistical performance claim follow from these single final runs.

The near-zero cases check a different property: a small but resolved gap must be distinguished from an unresolved or singular localizer. With the index acceptance threshold  $10^{-8}$ , a dense gap of  $10^{-6}$  is enclosed by  $[9.9994941 \times 10^{-7}, 1.0000076 \times 10^{-6}]$  and accepted. A gap of  $10^{-10}$  is bracketed by  $[5.8207661 \times 10^{-11}, 1.1641532 \times 10^{-10}]$  but receives no index; the exactly singular case also receives no index. For  $H = (\frac{1}{20} \frac{20}{399})$ ,  $S = 2I$ , and  $L = \text{diag}(H, -H)$ , the generalized gap is  $1.2499922 \times 10^{-3}$ , before and after a nonorthogonal basis change. The sparse interval is  $[1.2499921, 1.2499922] \times 10^{-3}$  at printed precision. This example tests the spectral pencil, not an estimate from LDL pivot magnitudes. All intervals are finite-precision numerical brackets, not interval-arithmetic certificates.

## S7.2 Sparse Pfaffian signs with a spin-mixed model

Time-reversed copies of the side-10 model are coupled by an onsite spin-mixing amplitude  $0.1 + 0.05i$  with the antisymmetric orbital structure required by physical time reversal. The complex localizer has order 800. The nontrivial case, a mass-shifted trivial case, and a complex nonorthogonal transformation of the nontrivial case give the results in Table S8. Both dense and sparse paths use the same trivial-reference orientation. MUMPS checks the metric and gap independently of the Tacho skew factorization.

Separate polynomial/permutation tests cover 338 dense/sparse Pfaffian comparisons of orders 2–24, scales from  $10^{-120}$  to  $10^{120}$ , disconnected blocks, and odd simultaneous row/column permutations. Deliberately singular cases emit zero-pivot diagnostics and are checked as rejected outcomes, not counted as successful topological classifications. Broken time reversal and invalid metrics are likewise rejected. The factorization itself remains on one rank; these tests do not demonstrate MPI-distributed Pfaffian factors or large SOC-material scaling.

| Case                          | $\nu$ | Gap (model units) | Solve-probe residual   |
|-------------------------------|-------|-------------------|------------------------|
| Spin-mixed nontrivial pair    | 1     | 0.79004778        | $2.28 \times 10^{-14}$ |
| Trivial pair                  | 0     | 1.1109451         | $2.54 \times 10^{-16}$ |
| Nonorthogonal nontrivial pair | 1     | 0.79004778        | $3.88 \times 10^{-15}$ |

Table S8: Dense/MUMPS/Tacho model comparison. Gap values agree at printed precision; the solve probes diagnose the sparse skew factors but are not rigorous global backward-error bounds.

### S7.3 Periodic AO localizer: mesh versus supercell

Ne is placed at  $(3, 3, 3)$  Å in a  $6 \times 6 \times 6$  Å cell periodic in XZ. The PBE/GPW calculation uses DZVP-MOLOPT-GGA-GTH-q8 from `BASIS_MOLOPT_UZH`, GTH-PBE-q8 SOC parameters, a single 2000 Ry grid, and an SCF tolerance of  $10^{-10}$ . A full Gamma-centered  $3 \times 1 \times 2$  SCF/property mesh is compared with a six-cell Gamma supercell and property mesh  $1 \times 1 \times 1$ . The query is  $(2.25, 3, 4.125)$  Å,  $E = -0.1$  hartree and  $\eta = 0.5$  hartree. Both spinor localizers have order 312 and give  $\nu = 0$ .

The generalized gap is 0.380994261391133 hartree for the primitive mesh and 0.380994261819095 hartree for the supercell, a difference of  $4.28 \times 10^{-10}$  hartree. The corresponding total energies are  $-34.923548219306682$  and  $-209.541289276241173$  hartree; normalizing the latter to one primitive cell leaves a difference of  $6.60 \times 10^{-9}$  hartree. This checks the full-torus Fourier construction with separately converged SCF potentials. It is neither an irreducible-property-mesh test nor a nontrivial topological material result. The high cutoff reduces grid error for this comparison; it does not establish a complete basis, size or  $\eta$  convergence study.

### S7.4 Molecular quadratic states and degenerate subspaces

Bi<sub>2</sub> is aligned along x with bond length 2.8 Å in an isolated 8 Å box. The PBE calculation uses TZVP-MOLOPT-GGA-GTH-q5, GTH SOC, a 150 Ry cutoff, and SCF tolerance  $10^{-9}$ . At the molecular center,  $E = -0.1$  hartree and  $\kappa = 0.1$  hartree/bohr, four states are requested in a 68-dimensional spinor AO space. The iterative trial space is bounded at 32 vectors, with a  $10^{-9}$  eigenpair-residual threshold. The GPW and GAPW pseudopotential inputs each compare dense and iterative solvers on four MPI ranks and two OpenMP threads.

The two Kramers pairs have quadratic gaps 0.319693242355779 and 0.320631445276286 hartree (the latter rounded across the pair). In both methods, the largest dense/iterative gap difference is  $6.0 \times 10^{-15}$  hartree. The iterative solve takes 145 Ritz iterations and its largest reported eigenpair residual is  $8.22 \times 10^{-10}$  hartree<sup>2</sup>, compared with  $2.82 \times 10^{-15}$  hartree<sup>2</sup> for the dense solution. Thus agreement of the eigenvalues alone would overstate the accuracy of the iterative eigenvectors.

Reanalysis of the stored coefficients verifies  $C^\dagger SC = I$  to  $1.1 \times 10^{-14}$  and all four singular values of  $C_{\text{dense}}^\dagger SC_{\text{iterative}}$  to within  $10^{-10}$  of unity. This compares the complete subspaces, not arbitrary individual phases or rotations within a Kramers pair. The reported decomposition  $\lambda = \delta E^2 + \kappa^2 \sum_j \delta x_j^2$  is consistent to  $2.9 \times 10^{-15}$  hartree<sup>2</sup>. These residuals refer to projected AO positions, not independently integrated continuum position variances. The small finite-cutoff calculation validates the solver and SOC operator integration; it is not a basis-converged molecular localization prediction.

### S7.5 Periodic quadratic coordinates and solver comparisons

The additional checks on 25 September 2026 use the paired trigonometric AO operators defined in the main text. The raw inputs, output logs, quadratic-state coefficients, source snapshot, and comparison script are retained in `validation/development/periodic-quadratic-records.tar.gz`. The accompanying `periodic-quadratic-summary.json` records the individual values and comparison bounds. The archive is small because it excludes restart wavefunctions and regenerable builds.

Unless specified otherwise, one He or Ne atom is placed at (2.25, 2.25, 2.25) Å in a cubic 4.5 Å cell. The scalar SCF uses LDA/PADE, a single 300 Ry grid, relative cutoff 40 Ry, and SCF tolerance  $10^{-10}$ . Basis functions are DZVP-MOLOPT-GGA-GTH-q2 for He and q8 for Ne from BASIS\_MOLOPT\_UZH; the GTH-PBE potentials, including Ne SOC parameters, come from GTH\_SOC\_POTENTIALS. These deliberately fixed small integration fixtures are not basis-converged predictions or recommended material settings. Both scalar and property meshes are Gamma-centered. The query is (1.875, 2.0, 2.625) Å,  $E = -0.1$  hartree, and  $\kappa = 0.1$  hartree/bohr. Two low quadratic states are requested, with iterative subspace limit 24, iteration limit 500, and eigenpair tolerance  $10^{-10}$  hartree<sup>2</sup> at these eigenvalues.

| Fixture     | Periodicity | Mesh                  | Lowest gap<br>(hartree) | Maximum gap difference<br>(hartree) |
|-------------|-------------|-----------------------|-------------------------|-------------------------------------|
| He, GPW     | Y           | $1 \times 2 \times 1$ | 0.358631854242405       | $1.0 \times 10^{-15}$               |
| He, GAPW    | XY          | $2 \times 1 \times 1$ | 0.351912947825943       | $< 10^{-15}$                        |
| He, GPW     | XYZ         | $2 \times 2 \times 1$ | 0.385841580050524       | $< 10^{-15}$                        |
| Ne, GPW/SOC | XZ          | $2 \times 1 \times 1$ | 0.406734402392864       | $5.0 \times 10^{-15}$               |

Table S9: Periodic quadratic dense/iterative comparisons. The gap difference is the maximum over the two requested states, comparing serial LAPACK with two-rank MPI/DBCSR/MUMPS. Each process uses two OpenMP threads and one OpenBLAS thread. Values use the recorded printed precision; vanishing differences do not imply exact arithmetic.

Including the reduced-mesh, wrapped-query and implicit-Gamma supercell variants gives 14 compared gaps, all within  $5.0 \times 10^{-15}$  hartree. The largest dense/iterative difference in the energy residual is  $1.94 \times 10^{-13}$  hartree. The iterative Ne/SOC eigenpair residual reaches  $6.67 \times 10^{-11}$  hartree<sup>2</sup>; it is reported separately from eigenvalue agreement and is not a bound on an individual eigenvector within a degenerate pair. No claim of gauge-matched periodic AO coefficients is inferred from these gap comparisons. Re-parsing all successful cases, including the additional cell and mesh variants, recovers the energy/chord decomposition of the squared gap within  $4.49 \times 10^{-15}$  hartree<sup>2</sup>.

For He, switching full-mesh SCF to atomic symmetry reduction or shifting the query by torus vectors to (10.875, -7.0, 7.125) Å changes the gaps by at most  $1.1 \times 10^{-15}$  hartree. The analysis still assembles the full torus; only the preceding SCF is symmetry reduced. Equivalent MONKHORST-PACK, unshifted MACDONALD and explicitly listed GENERAL SCF meshes also agree within  $1.1 \times 10^{-15}$  hartree.

An independently converged  $2 \times 2 \times 1$  He supercell, with property mesh  $1 \times 1 \times 1$ , gives gaps 0.385841580972906 and 0.391994927003925 hartree through the implicit-Gamma path. The largest difference from the primitive mesh is  $9.30 \times 10^{-10}$  hartree. The primitive SCF energy is -2.832198277992086 hartree; the supercell energy -11.328793115636264 hartree differs by  $9.17 \times 10^{-10}$  hartree after division by four. An explicit complex Gamma KPOINTS calculation agrees with the implicit-Gamma gaps within  $2.2 \times 10^{-12}$  hartree. The latter comparison checks the two input paths, not exact equality of separately converged SCF potentials or convergence in basis, mesh and cutoff.

Axis-permuted He chains cover X, Y and Z; GAPW slab variants cover XY and YZ, with Ne/SOC exercising XZ. A skew cell with rows (4.5, 0, 0), (0.6, 4.4, 0) and (0.3, 0.2, 4.6) Å is compared with the rigidly rotated cell, atom and query using

$$R = \begin{pmatrix} 0.6 & -0.8 & 0 \\ 0.8 & 0.6 & 0 \\ 0 & 0 & 1 \end{pmatrix}.$$

The resulting gaps 0.392476754147977 and 0.412631258413744 hartree are preserved within  $1.1 \times 10^{-15}$  hartree. This checks input-frame handling, not invariance under changing the primitive

basis of the reciprocal-coordinate localization metric. Separate invalid inputs are rejected for a mesh along an open direction, a mesh supplied with the finite formulation, an oversized AO torus, and inconsistent CELL/POISSON periodicity.

The low-level seven-observable tests independently check exact diagonal spectra, complex hopping, nonsingular AO basis changes, paired sine/cosine origin rotations and query wrapping. These algebraic tests supplement, rather than replace, the Gaussian operator comparisons. They also pass with bounds checking, undefined-behavior instrumentation and floating-point traps in the numerical kernel. The complete affected quadratic/localizer regression directories pass 52/52 serial and 87/87 two-rank MPI assertions, retaining the existing finite-system references. Neither the small torus tests nor the local two-rank runs establish large-system or multi-node scaling.

## **S8 Submission Validation Still Required**

The current record does not establish a nontrivial, size-converged DFT localizer classification for a large SOC flake, multi-node sparse scaling, or general 230-space-group TQC. The final archive should contain complete material inputs and potentials, basis and cutoff studies, spatial/energy/scale scans, independent band references, and timed runs with whole-process memory. Development diagnostics and planned work must not be merged into tables of completed material results.